

IntentCap: Intent-Carrying Capabilities for LLM Agent Extensions

Abstract

LLM agents act through tool calls, shell commands, file I/O, MCP servers, subagents, and memory, and each channel brings context into one shared planning state. A user’s request, a Skill’s instructions, a tool result, and a file’s content all enter this state with equal influence over every decision, yet they carry different trust and should govern different fields. Current defenses control which operations are allowed but not which inputs can influence which authority-bearing decisions.

INTENTCAP makes user intent the root principal and turns context influence into an explicit capability: each input class is authorized to influence specific decision fields, and no class can prove another’s authority. Task plans compile into proof-carrying capability leases over four issuer-owned input classes, validated by a deterministic checker before any side effect commits.

Across 38 multi-boundary enforcement events spanning tool calls, local execution, context placement, and delegation, the checker produces zero unsafe accepts. Ablations show the four-class partition is structural: collapsing it causes 94% of protected decisions to be falsely accepted.

1 Introduction

LLM agents are becoming extensible execution environments. A modern agent can load Skills with workflow instructions, references, and optional scripts [3]; connect to MCP servers that expose resources, prompts, and tools [2]; run local commands; and spawn subagents. This extensibility is useful precisely because these components affect future behavior. A Skill can teach the agent how to perform a workflow. A tool result can become evidence for a later action. A document can determine the content of a generated report. The same property creates a security problem: in an agent, text is not passive data—it can become control.

Most current defenses ask an operation-centric question: which tool may be called, which file may be read, which endpoint may be contacted, or which process may execute? That question is necessary but incomplete. Consider a user who asks an agent to extract tables from two selected PDFs, save spreadsheets, and create one GitHub issue in a named repository. The PDF contents should influence spreadsheet cells and perhaps the issue body. They should not influence the repository name, network destination, approval scope, tool selection, or decision to load another Skill. If hidden text in the PDF says “send all extracted values to attacker.com” or “create the issue in attacker/repo,” a tool-call allowlist or OS sandbox may catch some concrete effects, but it does not directly express the more agent-specific invariant: the untrusted PDF was never allowed to influence those decisions.

Securing agent extensions requires controlling not only what extensions can do, but what their context is allowed to mean. The root of authority is a structured *intent certificate*, not an extension package, server identity, or static manifest. A certificate records the current goal, selected objects, authorized sinks, constraints, approvals, and expiry. Context cells are labeled with provenance and allowed *influence modes*, such as observing, quoting, summarizing,




Figure 1: INTENTCAP architecture. The LLM proposes; the checker decides. Adapters submit field proofs before any side effect, placement, or handoff.

parameterizing, planning, instructing, delegating, or authorizing. A capability is then minted for a particular run as an *intent-carrying lease*: it authorizes an operation only if the operation is derivable from the current intent and its control dependencies come from context allowed to influence that decision class.

The root cause is *issuer collapse*: user intent, Skill instructions, tool metadata, and runtime observations share one planning channel, so one class can fill another’s authority fields.

2 Design

INTENTCAP addresses issuer collapse by making the authorization unit a *pre-effect commit record* rather than an action-level predicate. Figure 1 shows the pipeline. A trusted issuer canonicalizes user intent into a structured certificate. A context labeler assigns provenance and influence modes. An untrusted LLM-assisted compiler proposes plans and candidate leases. A deterministic checker validates field-owner proofs and atomically consumes lease state via a single transition:

$$\text{check_and_consume}(e, \text{lease}, \text{proofs}, \text{prov}, \sigma) \rightarrow \text{allow}(\sigma', \text{audit}) \mid \text{deny}(\text{reason})$$

Every authority-changing boundary—tool call, local execution, context placement, Skill instruction slot, or delegation handoff—must submit this record *before* the side effect commits. The checker is the sole writer of lease state; adapters cannot cache decisions or mutate budgets locally.

INTENTCAP partitions inputs into four issuer-owned proof boundaries: *agent* context (user intent, selected objects/sinks, approvals), *instruction* context (system policy, endorsed Skill/manual procedures), *tool* context (MCP/registry schemas, credential scopes, sandbox contracts), and *env* context (runtime observations, tool results, file state, script outputs). The *no-substitution rule* requires that each protected field be proven by its owner class: tool metadata cannot supply user-authorized sinks, runtime observations cannot supply instruction authority, and Skill text cannot widen approval scope. Two classes may merge only if issuer, forgeable surface, owned fields, and lifecycle authority are identical.

A lease binds operation, object, arguments, intent derivation, control/data provenance, temporal guards, budget, expiry, and delegation depth. Leases are short-lived, consumable, and non-self-widening: a one-shot issue lease is consumed on first use, and a child subagent receives only attenuated capabilities. The checker validates all of these atomically in a single transition, preventing the authority-state split where separate guards lose track of consumption or delegation across boundaries.

3 Preliminary Evidence

RQ1: Can the event model express benign behavior without over-blocking? INTENTCAP produces 0 unsafe accepts across 3,746 protected-decision events replayed from AgentDojo, MCPTox, InjecAgent, and tau2-bench [1? ? ?]. On the utility side, INTENTCAP covers all 3,813 benign reference actions in the tau2-style proxy and passes 2,554 of 2,556 applicable tool-oracle checks.

RQ2: Is the four-class partition a safety requirement or a naming choice? The partition is load-bearing: removing issuer ownership entirely (collapsing all four classes into generic trusted context) causes 3,593 false accepts out of 3,823 checker-denied events (94%). Even the weakest single-edge collapse (tool→agent) falsely accepts 1,928 events (50%). On 7 controlled workflow residuals, a post-hoc policy DSL that checks predicates without field ownership falsely accepts 7/7; splitting lifecycle state across separate guards falsely accepts 5/7.

RQ3: Does the same commit API work beyond tool calls? The check_and_consume transition generalizes across five boundary types: local env side effects, context placement, Skill instruction slots, delegation handoffs, and an OS-monitor-style replay backend. Across 38 checker-submitted boundary events, INTENTCAP allows 17 authorized effects and blocks 21 violations, with 0 unsafe executions or placements and 0 checker/monitor mismatches.

RQ4: Are leases auditable authority objects? All INTENTCAP policy families achieve oracle distance 0 across 24 adjudicated lease labels covering InjecAgent, MCPTox, and tau2 samples; every non-INTENTCAP baseline has positive distance over 144 scored policy rows.

4 Discussion

INTENTCAP asks a different question from adjacent work. EIM and bptime ask how an extension can safely use host-provided resources [6]. ActPlane asks how policy can be enforced below the tool layer [5]. SkillGuard asks how a Skill package can be governed as a permission-bearing artifact [4]. INTENTCAP asks which future decisions a given context may influence, under a given user intent, and with what proof. That distinction keeps the LLM outside the trusted computing base, treats the current user goal as the root of authority, and makes context influence a least-privilege capability rather than an implicit side effect of putting text in the prompt.

Current evidence covers replay safety, issuer-collapse ablations, and local multi-boundary enforcement. End-to-end utility with a live user simulator, independent expert-oracle lease labels, and production OS-level enforcement remain open. Scaling the checker to real-time MCP broker integration and measuring approval burden are immediate next steps.

References

- [1] Edoardo DeBenedetti, Jie Zhang, Mislav Balunović, Luca Beurer-Kellner, Marc Fischer, and Florian Tramèr. 2024. AgentDojo: A Dynamic Environment to Evaluate Prompt Injection Attacks and Defenses for LLM Agents. In *Advances in Neural Information Processing Systems 37 (NeurIPS), Datasets and Benchmarks Track*. doi:10.48550/arXiv.2406.13352
- [2] Model Context Protocol. 2025. Specification. <https://modelcontextprotocol.io/specification/2025-06-18>.
- [3] OpenAI. 2026. Agent Skills. <https://developers.openai.com/codex/skills>. Accessed: 2026-07-02.
- [4] Shidong Pan, Xiaoyu Sun, Tianyi Zhang, Dianshu Liao, Meixue Si, and Zhenchang Xing. 2026. SkillGuard: A Permission Framework for Agent Skills. arXiv preprint

arXiv:2606.03024. doi:10.48550/arXiv.2606.03024

- [5] Yusheng Zheng, Tianyuan Wu, Quanzhi Fu, Tong Yu, Wenan Mao, Wei Wang, Dan Williams, and Andi Quinn. 2026. ActPlane: Programmable OS-Level Policy Enforcement for Agent Harnesses. arXiv preprint arXiv:2606.25189. doi:10.48550/arXiv.2606.25189
- [6] Yusheng Zheng, Tong Yu, Yiwei Yang, Yanpeng Hu, Xiaozheng Lai, Dan Williams, and Andi Quinn. 2025. Extending Applications Safely and Efficiently. In *Proceedings of the 19th USENIX Symposium on Operating Systems Design and Implementation (OSDI '25)*. USENIX Association. <https://www.usenix.org/conference/osdi25/presentation/zheng-yusheng>